

# Proof of Concept Portfolio for Predictive Maintenance using TimesFM

## Key Objectives

I began by exploring how Large Language Models (LLMs) could be applied to predictive maintenance in real-world industrial setups. Initially, I studied the traditional approaches to predictive maintenance and investigated the direction in which recent AI developments are heading.

It became clear that general-purpose LLMs like ChatGPT or LLaMA aren't directly suitable for handling time series data. I then shifted focus to transformer-based architectures, which form the backbone of LLMs, and identified **TimesFM**—a foundation model pretrained by Google specifically on time series data—as a strong candidate for this task.

## Dataset for Analysis

To begin, I chose the **NASA Bearing dataset** because it presents raw signals in a relatively simple form. This helped me understand fundamental concepts such as signal windowing, feature engineering, and degradation modeling.

Using this dataset, I built a **custom health index** from features like mean, standard deviation, and kurtosis, and mapped it to Remaining Useful Life (RUL) using threshold-based logic.

However, for more robust modeling and demonstration, I plan to use the **CMAPSS Turbofan dataset**, which provides structured labels and standardized RUL calculations—allowing me to focus more on model design and less on preprocessing challenges.

## What I've Built So Far

### 1. Data Preprocessing on NASA Bearing Data

- Developed a pipeline to extract sliding-window features (mean, std, kurtosis) from raw vibration signals.
- Engineered timestamps and a synthetic health index to simulate degradation.
- Computed RUL dynamically by associating the health index with a failure threshold (RUL = 0 when health index crosses it).
- Simulated various degradation patterns and visualized them.

### 2. Forecasting with TimesFM (Zero-Shot)

- Used **TimesFM in inference mode** (no fine-tuning) to forecast future values of the custom health index.
- Estimated RUL from the forecasted trajectory.
- Demonstrated TimesFM's ability to operate without training on the same asset, showcasing its foundation model strength.

### 3. TimesFM Fine-Tuning

- Adapted the official TimesFM finetuning framework with enhancements:
  - **Distributed training support** via PyTorch DDP.
  - **Quantile loss** for capturing uncertainty in forecasts.
  - **WandB integration** for tracking experiments.
- Prepared patch-level data formatting to align with TimesFM's input-output specification.

### 4. Model Evaluation Strategy

- Built visual diagnostics to track degradation and predicted RUL over time.
- Compared model forecasts with baseline trends.
- Evaluated using standard metrics such as **MSE**.

## What's Ready for Full Implementation

- Modular data loader utilities for sliding windows and patch alignment.
- Preprocessing pipeline for CMAPSS or similar datasets.
- Working zero-shot inference with TimesFM.
- Fine-tuning scaffold that supports custom loss functions and multi-GPU training.
- Evaluation framework to compare predicted vs. actual RUL using RMSE, MAE, and custom degradation metrics.

## Tools & Technologies Used

- **LLM for Time Series:** TimesFM (via Hugging Face), PyTorch

- **Preprocessing & Modeling:** pandas, numpy, scipy, sklearn
- **Visualization:** matplotlib, seaborn
- **Experiment Logging:** Weights & Biases (WandB)
- **Training Infrastructure:** PyTorch DDP (Distributed), Custom Fine-Tuner
- **Problem Focus:** RUL Estimation, Health Index Forecasting, Failure Prediction

## Next Steps

- I can wrap up the POC phase and present the results based on current NASA bearing data.
- Or, I can extend this work to a **full implementation on the CMAPSS dataset**, which better mirrors real-world predictive maintenance settings and can be a compelling demonstration of TimesFM's capabilities.